

# Possible `ergm` bug involving `ergm.egrad`

December 17, 2017

The function `ergm.egrad` of R package `ergm` evaluates the gradient of the natural parameter vector. For geometrically weighted model terms, the gradient may be incorrect. The following is based on R package `ergm` version 3.8.0 from 2017-08-18, but R package `ergm` version 2.0-1 from 2008-01-19 uses the same gradient as R package `ergm` version 3.8.0, as downloaded from <https://cran.r-project.org/src/contrib/Archive/ergm/>.

## 1 The gradient of the natural parameter vector

The natural parameters are defined as

$$\eta_i(\boldsymbol{\theta}) = \theta_1 \exp(\theta_2) (1 - (1 - \exp(-\theta_2))^i), \quad i = 1, \dots, n,$$

so that the partial derivatives are given by

$$\begin{aligned} \nabla_{\theta_2} \eta_i(\boldsymbol{\theta}) &= \theta_1 [\exp(\theta_2) (1 - (1 - \exp(-\theta_2))^i) - \exp(\theta_2) i (1 - \exp(-\theta_2))^{i-1} \exp(-\theta_2)] \\ &= \theta_1 [\exp(\theta_2) (1 - (1 - \exp(-\theta_2))^i) - i (1 - \exp(-\theta_2))^{i-1}]. \end{aligned}$$

## 2 Implementation of gradient in R package `ergm`

The R script

```
library(ergm)
net <- network.initialize(5)
model <- ergm.getmodel(net ~ gwesp(fixed = FALSE), net)
model$etamap$curved[[1]]$map
model$etamap$curved[[1]]$gradient
```

shows

```
> model$etamap$curved[[1]]$map
function (x, n, ...) {
  i <- 1:n
  x[1] * exp(x[2]) * (1 - (1 - exp(-x[2]))^i)
}
```

and

```
> model$etamap$curved[[1]]$gradient
function (x, n, ...) {
  i <- 1:n
  a <- 1 - exp(-x[2])
  exp(x[2]) * rbind(1 - a^i, x[1] * (1 - a^i - i * a^(i - 1)))
}
```

The implementation of the gradient is equivalent to

$$\nabla_{\theta_2} \eta_i(\boldsymbol{\theta}) = \theta_1 [\exp(\theta_2) (1 - (1 - \exp(-\theta_2))^i) - \mathbf{exp}(\boldsymbol{\theta}_2) i (1 - \exp(-\theta_2))^{i-1}].$$

The boldfaced  $\mathbf{exp}(\boldsymbol{\theta}_2)$  should not be here.

### 3 The correct gradient

The correct gradient is given by

```
gradient <- function(x, n) {
  i <- 1:n
  a <- 1 - exp(-x[2])
  rbind(exp(x[2]) * (1 - a^i), x[1] * (exp(x[2]) * (1 - a^i) - i * a^(i - 1)))
}
```

## 4 Example

A simple example is given by

$$\eta_1(\boldsymbol{\theta}) = \theta_1 \exp(\theta_2) (1 - (1 - \exp(-\theta_2))) = \theta_1,$$

which implies that

$$\nabla_{\theta_2} \eta_1(\boldsymbol{\theta}) = 0.$$

R package `ergm` reports

```
> model$etamap$curved[[1]]$gradient(c(1,0.5), 3)
      [,1]      [,2]      [,3]
[1,]  1.0000000  1.3934693  1.5482875
[2,] -0.6487213  0.0960268  0.7825317
```

which shows that the  $\nabla_{\theta_2} \eta_1(\boldsymbol{\theta})$  is evaluated as  $-0.6487213$  rather than 0.

The correct gradient

```
gradient <- function(x, n) {
  i <- 1:n
  a <- 1 - exp(-x[2])
  rbind(exp(x[2]) * (1 - a^i), x[1] * (exp(x[2]) * (1 - a^i) - i * a^(i - 1)))
}
```

reports

```
> gradient(c(1,0.5), 3)
      [,1]      [,2]      [,3]
[1,]    1  1.3934693  1.548287
[2,]    0  0.6065307  1.083833
```

which shows that  $\nabla_{\theta_2} \eta_1(\boldsymbol{\theta})$  is evaluated as 0, which is indeed correct.

## 5 Possible implications

If  $\theta_2 > 0$ , then  $\exp(\theta_2) > 1$ . Therefore, R package `ergm` subtracts from all partial derivatives a term that is too large, because it is multiplied by  $\exp(\theta_2) > 1$ . An example is  $\nabla_{\theta_2} \eta_1(\boldsymbol{\theta})$ , which

is 0 whereas R package `ergm` reports  $-0.6487213$ . Hence small, positive partial derivatives can turn into negative partial derivatives, so that the signs of the partial derivatives are wrong and the Monte Carlo maximum likelihood algorithm implemented in `ergm` may venture into the wrong direction. It is possible that this can lead to convergence issues.